

Optimization for large scale machine learning

ENS, M2 Info

Mathurin Massias
mathurin.massias@gmail.com

Contents

1	Motivation: gradient descent on ordinary least squares	3
1.1	Why does statistical learning need optimization?	3
1.2	Solving least squares with gradient descent	5
1.3	Numerical puzzles	7
1.4	Exercises	8
2	Reminder on convex analysis	9
2.1	Convexity and minimizers	9
2.2	Exercises	10
2.2.1	Convexity	10
2.2.2	Gradient	11
3	Gradient descent on convex functions	12
3.1	Exercises	16
3.1.1	Convexity inequalities	16
3.1.2	Gradient descent	16
4	Non smooth convex optimization	18
4.1	Constrained optimization	18
4.2	Extended value functions	19
4.3	Subdifferential of convex functions	19
4.4	The proximal operator	21
4.5	The proximal point algorithm	22
4.6	The proximal gradient descent algorithm	23
4.7	Exercises	24
4.7.1	Subdifferential	24
4.7.2	Constrained optimization	24
5	Duality	25
5.1	The Fenchel transform	25
5.2	Fenchel-Rockafellar duality	26

5.3	Primal dual algorithms	28
5.4	Exercises	30
6	Second order optimization: Newton method	31
6.1	Newton method	31
6.2	Quasi Newton methods	33
6.2.1	The Broyden/SR1 formula	33
6.2.2	DFP and BFGS	34
6.3	Exercises	34
7	Continuous view: gradient flows	35
8	Optimal first order methods: Nesterov acceleration	35
9	The blessing of non-smoothness: Manifold identification of proximal algorithms	35
10	Fixed point schemes	35
11	Lower bounds	35
12	Stochastic methods: SGD	35
13	Stochastic methods: Variance reduced methods	35
14	Lab3: stochastic methods + DL	35
15	Non convex optimization: sparsity	35
16	Optimization for Deep Learning	35
17	Mirror descent	36

Goals of the class

In this class we will study the resolution of optimization problems arising in Machine Learning. These problems are usually characterized by their large scale, which call for dedicated classes of algorithms. They also usually present some specific structure, that can be leveraged.

The main objectives are to:

1. develop a taxonomy of optimization problems and which algorithms can be used in which settings
2. understand convergence properties and functions properties that govern them
3. master theoretical tools and techniques to derive convergence proofs
4. get hands on practice to challenge theoretical results, in Python

I thank Cesare Molinari, Saverio Salzo and Silvia Villa for their graduate course on convex analysis.

General notational conventions

We almost always work in finite dimension and the optimization is performed over the set \mathbb{R}^d . In a Machine Learning context, the integer d will thus denote the dimension of the parameter space, while n will denote the number of samples (observed data points). We adopt the optimization convention, i.e. we write a linear system as $Ax = b$ (while statistics has $y = X\beta$, the inverse problems field writes $f = Au$, signal processing writes $y = \Phi x \dots$): the unknown is x ; the iterates are x_k or x_t . As much as possible we write x^* for the minimizer of given problems, but due to the various meaning of star (Fenchel conjugation, adjoint), it happens that we also write \hat{x} as in the statistical literature.

These conventions are meant to ease reading, but there can always be exceptions. It is anyways a good exercise to adapt to different notation.

1 Motivation: gradient descent on ordinary least squares

Most optimization problems studied in this class are of the following form:

$$\inf_{x \in \mathcal{C}} f(x) \quad , \quad (1.1)$$

where \mathcal{C} is a subset of \mathbb{R}^d and f is a real-valued function on a subset of \mathbb{R}^d .

Formally, solving [Problem \(1\)](#) means to find the largest lower bound on all values $f(x)$ when x describes \mathcal{C} . This largest lower bound, call the *infimum*, may not exist, and if it exists it may not necessarily be attained, even for nicely-behaved f . The reader should mind that most of the time, we write $\min_{\mathcal{C}} f$ instead of $\inf_{\mathcal{C}} f$; this is an abuse of notation and does not mean that the infimum is attained (not even that it is finite).

In Machine Learning, rather than the smallest value taken by f , we are more interested in finding a minimizer, that is

$$x^* \in \operatorname{argmin}_{x \in \mathcal{C}} f(x) \quad , \quad (1.2)$$

meaning in finding $x^* \in \mathcal{C}$ such that for all $x \in \mathcal{C}$, $f(x^*) \leq f(x)$. In Machine Learning, x^* usually represent the parameters of a model, that once computed are used to achieve some statistical task.

1.1 Why does statistical learning need optimization?

Let $n \in \mathbb{N}$. Suppose that we have collected a dataset $(a_i, b_i)_{i \in [n]}$ of pairs observations-targets, and we assume that they are linearly related, i.e. there exists $x \in \mathbb{R}^d$ such that $b_i \approx a_i^\top x$.

This is very vague. To make it more precise and rigorous, it is extremely frequent to model the dataset as n independent, identically distributed realizations of two random variables A and B linked by

$$B \sim \mathcal{N}(A^\top x_{\text{true}}, \sigma^2) \quad , \quad (1.3)$$

meaning that for all $i \in [n]$, $b_i = x_{\text{true}}^\top a_i + \varepsilon_i$, with $(\varepsilon_i)_{i \in [n]}$ n realizations of a centered Gaussian variable of variance σ^2 .

Given the collected observations, how can we infer x_{true} ? From a statistical point of view, an answer is to maximize the conditional likelihood. Under model (1.3), the conditional likelihood is:

$$\prod_{i=1}^n \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{\|a_i^\top x - b_i\|^2}{2\sigma^2}\right) . \quad (1.4)$$

Minimization of the negative log likelihood in x gives

$$x^* \in \operatorname{argmin}_{x \in \mathbb{R}^d} \sum_{i=1}^n (b_i - a_i^\top x)^2 = \operatorname{argmin} \frac{1}{2} \|Ax - b\|^2 , \quad (1.5)$$

where the **design matrix** $A \in \mathbb{R}^{n \times d}$ has i -th row a_i , and the **target vector** $b \in \mathbb{R}^n$ has i -th entry b_i . [Problem \(5\)](#) is called Ordinary Least Squares (OLS). Solving it gives an estimator for x_{true} and allows, for example, predicting a value b_{n+1} if one observes a new entry point a_{n+1} . This is also an illustration that the value of the infimum is not always very interesting (if $d \geq n$ and A has rank n it is 0), contrary to the minimizer.

How to compute the solution of [Problem \(5\)](#)? This is a classical problem in linear algebra and there exist direct ways to do it ([Exercise 1.6](#)). But in other settings there may not exist a closed-form solution: it is perfectly possible to postulate a different generative model than (1.3), or simply not impose a generative model, and find x by minimizing another cost/loss function. This is the setting of empirical risk minimization (ERM): one finds the parameters of a model by minimizing a cost function, often –but not always– arising from likelihood maximization.

Example 1.1 (ℓ_1 and Huber regression). *Ordinary least squares are strongly influenced by outliers: the squared loss gives a lot of importance to large differences. In the objective function (5) a mismatch of 2 between the true observation b_i and the model fit $a_i^\top x$ costs $2^2/2 = 2$, but a mismatch of 10 costs $10^2/2 = 50$. A more robust solution, that is less dominated by large mismatches (or “outliers”), can be found by minimizing the sum of absolute errors:*

$$\operatorname{argmin}_{x \in \mathbb{R}^d} \sum_{i=1}^n |a_i^\top x - b_i| = \operatorname{argmin}_{x \in \mathbb{R}^d} \|Ax - b\|_1 . \quad (1.6)$$

[Problem \(6\)](#) is unfortunately more complicated to solve, as we shall see. It can be reformulated as a Linear Program, but the cost to solve it scales very badly with d . A hybrid of [Problems \(5\)](#) and [\(6\)](#) involves the Huber loss ([Figure 1](#)):

$$\operatorname{argmin}_{x \in \mathbb{R}^d} \sum_{i=1}^n \text{huber}_\rho(a_i^\top x - b_i) , \quad (1.7)$$

with the Huber loss defined as:

$$\text{huber}_\rho(x) = \begin{cases} \frac{x^2}{2\rho} + \frac{\rho}{2} , & \text{if } |x| \leq \rho , \\ |x| , & \text{otherwise.} \end{cases} \quad (1.8)$$

The Huber loss is less sensitive to outliers than the squared ℓ_2 loss, and more optimization friendly than the ℓ_1 loss, but there is no free lunch: it relies on a parameter $\rho \in \mathbb{R}_+$ that needs to be tuned, and there is no obvious and cheap way to choose the best ρ .

For both problems, there is no closed-form solution and we need optimization techniques to approximate the minimizer.




Figure 1: ℓ_1 , squared ℓ_2 and Huber ($\rho = 1$) losses

Finally, instead of a linear model, one may model the dependency between observation and target by a function belonging to some parametric class $\{f_x : \mathbb{R}^d \mapsto \mathbb{R} : x \in \mathbb{R}^p\}$, described by p parameters stored in x :

$$b_i \approx f_x(a_i) \quad , \quad (1.9)$$

and, depending on what is meant by \approx , minimize a loss $L : \mathbb{R} \times \mathbb{R} \rightarrow \mathbb{R}$, yielding the following Empirical Risk Minimization (ERM) problem:

$$x^* \in \operatorname{argmin}_{x \in \mathbb{R}^p} \sum_{i=1}^n L(f_x(a_i), b_i) \quad . \quad (1.10)$$

As we have seen, the typical example is $f_x = x^\top \cdot$; but this framework also fits the Deep Learning framework (with f_x a composition of matrix vector multiplications followed by pointwise application of a non-linear function.).

1.2 Solving least squares with gradient descent

Problem (5) can be solved in closed-form with dedicated algorithms such as Cholesky, LU or QR decompositions (see [Exercise 1.6](#)); but ℓ_1 , Huber and generic ERM models cannot. We need general purpose algorithms to solve classes of similar problems.

One of the arguably simplest minimization algorithm is gradient descent (GD). If it exists, the gradient of a function points in the direction of maximal increase, since the directional derivative $f'(x; v) = \lim_{\epsilon \rightarrow 0} \frac{f(x+\epsilon v) - f(x)}{\epsilon}$ is given by

$$f'(x; v) = \langle \nabla f(x), v \rangle \quad . \quad (1.11)$$

To decrease the value of f as fast as possible, it makes sense to move in the opposite direction of the gradient. Taking a stepsize $\eta > 0$, this gives the gradient descent iterations:

$$\begin{aligned} x_0 &\in \mathbb{R}^d, \\ x_{t+1} &= x_t - \eta \nabla f(x_t). \end{aligned} \quad (1.12)$$

Proposition 1.2 (Convergence rate of GD on OLS). *Consider the OLS problem, and assume that there exist strictly positive scalars μ and L such that $\mu \text{Id} \preceq A^\top A \preceq L \text{Id}$. Let $\kappa = L/\mu$ denote the condition number. Then there exists a unique minimizer of least squares, x^* , and the iterates of gradient descent on ordinary least squares with stepsize $\eta = 1/L$ satisfy:*

$$\|x_t - x^*\| \leq \exp(-t/\kappa) \|x_0 - x^*\|. \quad (1.13)$$

Proof. The gradient of $x \mapsto \frac{1}{2} \|Ax - b\|^2$ is $A^\top(Ax - b)$, hence gradient descent on ordinary least squares read

$$x_{t+1} = x_t - \eta A^\top(Ax_t - b). \quad (1.14)$$

There exists a unique minimizer x^* , and it satisfies $A^\top Ax^* = A^\top b$ ([Exercises 1.5](#) and [1.6](#)), so:

$$x_{t+1} - x^* = (\text{Id} - \eta A^\top A)(x_t - x^*). \quad (1.15)$$

Observe that

$$(1 - \eta L) \text{Id} \preceq \text{Id} - \eta A^\top A \preceq (1 - \eta \mu) \text{Id}, \quad (1.16)$$

hence

$$\|\text{Id} - \eta A^\top A\|_2 \leq q(\eta) \triangleq \max(|1 - \eta L|, |1 - \eta \mu|). \quad (1.17)$$

Setting $\eta = 1/L$ yields $q(\eta) = 1 - \mu/L$, and using $(1 - u)^t \leq \exp(-tu)$ concludes. \square

Remark 1.3. *The convergence rate (1.13), is a very good rate in the sense that it decays fast towards 0. It is called a linear rate¹. Suppose that we want to get ε -close to x^* , then we need at most $\kappa \log(\|x_0 - x^*\|/\varepsilon)$ iterations of GD, which grows quite slowly as ε goes to 0.*

We can also obtain a rate in objective value:

$$f(x_t) - f(x^*) = \langle \nabla f(x^*), x_t - x^* \rangle + \frac{1}{2} (x_t - x^*)^\top A^\top A (x_t - x^*) \quad (1.18)$$

$$= \frac{1}{2} (x_t - x^*)^\top A^\top A (x_t - x^*) \quad (1.19)$$

$$\leq \frac{L}{2} \|x_t - x^*\|^2, \quad (1.20)$$

but as we've seen, such rates are usually less desirable than rates on x_t .

Remark 1.4 (Other choices of η). *We have actually proven the stronger result:*

$$\|x_t - x^*\| \leq q(\eta)^t \|x_0 - x^*\|. \quad (1.21)$$

This shows that any choice of η in $]0, 2/L[$ leads to convergence, because it results in $q(\eta) < 1$. We can tune η to minimize $q(\eta)$: doing so yields $\eta = \frac{2}{L+\mu}$ and $q(\eta) = \frac{L-\mu}{L+\mu} = \frac{\kappa-1}{\kappa+1}$.

¹this is a terrible name: it is linear in log scale

Questions: Does there always exist an L as in the assumptions of [Proposition 1.2](#)? And a μ ? To what do they correspond geometrically?

1.3 Numerical puzzles

On [Figure 2](#) is the numerically observed convergence rate of GD on one instance of OLS. The matrix A is 500 by 500 with i.i.d. normal entries so it is full rank, $\mu > 0$ and [Proposition 1.2](#) applies. Yet the rate we see is clearly not linear. Why?




Figure 2: Objective convergence rate of GD on OLS on a random i.i.d. $A \in \mathbb{R}^{500 \times 500}$.

On [Figure 3](#), A is still random but its shape is 200×300 , so $A^\top A$ is not full rank, $\mu = 0$. Yet we see a clear linear convergence rate (numerical floating point errors kick in at iteration 1000). Why?




Figure 3: Objective convergence rate of GD on OLS on a random i.i.d. $A \in \mathbb{R}^{200 \times 300}$.

1.4 Exercises

Exercise 1.5. ☹ Let $A \in \mathbb{R}^{n \times d}$. Show that $\text{Ker } A = \text{Ker } A^*A$.

Show that for any $b \in \mathbb{R}^n$ there exist a solution to $A^*Ax = A^*b$.

☹☹ Show that there does not always exist a solution to $A^*Ax = A^*b$ in the infinite dimensional case (when A is a bounded linear operator between infinite dimensional Hilbert spaces.)

Exercise 1.6. ☹ Let $A \in \mathbb{R}^{n \times d}$, $b \in \mathbb{R}^n$. Show that solving Ordinary Least Squares:

$$\min \frac{1}{2} \|Ax - b\|^2, \quad (1.22)$$

amounts to solving $A^*Ax = A^*b$ (aka the normal equations).

Show that the set of solutions is:

$$A^\dagger b + \text{Ker } A.$$

Exercise 1.7. ☹ When is $x \mapsto \|Ax - b\|^2$ strictly convex? Strongly convex?

Exercise 1.8 (Least squares with intercept). ☹☹ An intercept x_0 is a constant scalar term in the linear prediction function, that becomes $a \mapsto a^\top x + x_0$. Fitting an intercept can be done by adding a column of 1s to A . Alternatively, show that the solution of least squares with intercept,

$$(\hat{x}, \hat{x}_0) \in \underset{x \in \mathbb{R}^d, x_0 \in \mathbb{R}}{\text{argmin}} \frac{1}{2} \|Ax - b - x_0 \mathbf{1}\|^2 \quad (1.23)$$

is given by:

$$\hat{x} = \hat{x}_c, \quad (1.24)$$

$$\hat{x}_0 = \frac{1}{n} \sum_1^n (a_i^\top \hat{x} - b_i), \quad (1.25)$$

where \hat{x}_c is the solution of least squares without intercept on centered data A_c and b_c (versions of A and b where the rowwise mean has been subtracted).

Exercise 1.9 (Gradient descent on isotropic parabola). ☹ Let $A \in \mathbb{R}^{n \times d}$ be such that the condition number² of $A^\top A$ is equal to 1. Show that gradient descent with stepsize $1/L$ converges in a single iteration for the problem $\min \frac{1}{2} \|Ax - b\|^2$.

²i.e. the ratio between the largest and the smallest eigenvalues of $A^\top A$.

2 Reminder on convex analysis

Section 1 has shown us basic notions in optimization. We have studied gradient descent, an iterative algorithm producing a sequence of iterates x_k that aim at solving an optimization problem. For the Ordinary Least Squares objective function, we have proved two types of convergence results³:

- in iterates: $\|x_k - x^*\| \rightarrow 0$
- in function values: $f(x_k) - \inf f \rightarrow 0$.

Both result were *non-asymptotic*: we had information about the speed at which convergence happened.

In the sequel we want to minimize functions beyond least squares. For starters, we'll work with convex functions, because they're roughly the only ones we can hope to minimize globally.

2.1 Convexity and minimizers

Definition 2.1 (Global and local minimizers). *Let $f : \mathbb{R}^d \rightarrow \mathbb{R}$. A global minimizer of f is a point x^* such that $f(x^*) \leq f(x)$ for all $x \in \mathbb{R}^d$. A local minimizer of f is a point x^* such that there exists a neighborhood V of x^* such that $f(x^*) \leq f(x)$ for all $x \in V$.*

Definition 2.2 (Convexity). *A function $f : \mathbb{R}^d \rightarrow \mathbb{R}$ is convex if and only if it lies below its cords:*

$$\forall x, y \in \mathbb{R}^d, \forall \lambda \in [0, 1], \quad f(\lambda x + (1 - \lambda)y) \leq \lambda f(x) + (1 - \lambda)f(y) . \quad (2.1)$$

It is strictly convex if the inequality (2.1) holds strictly (for $\lambda \in]0, 1[$).

Definition 2.3 (Strong convexity). *Let $\mu > 0$. A function $f : \mathbb{R}^d \rightarrow \mathbb{R}$ is μ -strongly convex if for all $x, y \in \mathbb{R}^d$,*

$$f(\lambda x + (1 - \lambda)y) \leq \lambda f(x) + (1 - \lambda)f(y) - \frac{\mu}{2}\lambda(1 - \lambda)\|x - y\|^2 . \quad (2.2)$$

Note: this is the only true definition of strong convexity. Other definitions are rather characterizations in special cases. In particular, many other definitions assume that the norm is the Euclidean one, a requirement that this definition does not have. This plays a role in Section 17.

When $\|\cdot\|$ is the Euclidean norm, then f is μ -strongly convex if and only if $f - \frac{\mu}{2}\|\cdot\|^2$ is convex (Exercise 2.16). In this sense, a strongly convex function is so convex that when you subtract a parabola, it remains convex.

Convex functions are amenable to optimization because of the nice “local to global” properties they enjoy.

Proposition 2.4 (Local minimizers are global minimizers). *Let $f : \mathbb{R}^d \rightarrow \mathbb{R}$ be a convex function. Then all local minimizers of f are also global.*

³Does one imply the other? Under which condition?

Proposition 2.5 (Above its tangents). *Let $f : \mathbb{R}^d \rightarrow \mathbb{R}$ be a convex differentiable function. Then*

$$\forall x, y \in \mathbb{R}^d, \quad f(x) - f(y) \leq \langle \nabla f(x), x - y \rangle .$$

Written as $f(y) \geq f(x) + \langle \nabla f(x), y - x \rangle$, it tells that f lies above all of its tangents. From local information $\nabla f(x)$, we get an information about the whole behavior of f .

Convex differentiable functions are nice because it is easy to characterize their minimizers.

Proposition 2.6 (Optimality condition for convex differentiable functions). *Let f be a convex differentiable function. Then*

$$x^* \in \operatorname{argmin}_{x \in \mathbb{R}^d} f(x) \Leftrightarrow \nabla f(x^*) = 0 . \quad (2.3)$$

Equipped with this theoretical framework, we can move to our first convergence proofs on generic functions.

2.2 Exercises

2.2.1 Convexity

Exercise 2.7. 🍷 *Show that local minimizers of convex functions are global minimizers.*

Exercise 2.8 (Pointwise supremum preserves convexity). 🍷 *Let $(f_i)_I$ be a family of convex functions (not necessarily countable). Show that $x \mapsto \sup_{i \in I} f_i(x)$ is convex.*

Exercise 2.9 (Precomposition by linear operator preserves convexity). 🍷 *Let $f : \mathcal{Y} \rightarrow \mathbb{R}$ be a convex function and $A : \mathcal{X} \rightarrow \mathcal{Y}$ a linear operator. Show that $f(A \cdot)$ is convex (on \mathcal{X}).*

Exercise 2.10 (Misconceptions on existence of minimizers). 🍷 *Provide an example of convex function which does not admit a minimizer.*

What if the function is continuous and lower bounded?

Exercise 2.11. 🍷 *Show that a strictly convex function has at most one minimizer.*

Exercise 2.12 (Jensen's inequality). 🍷 *Let $f : \mathbb{R}^d \rightarrow \mathbb{R}$ be convex. Let $n \in \mathbb{N}$, $x_1, \dots, x_n \in \mathbb{R}^d$, and let $\lambda_1, \dots, \lambda_n$ be positive scalars summing to 1. Show that $f(\sum_{i=1}^n \lambda_i x_i) \leq \sum_{i=1}^n \lambda_i f(x_i)$.*

Exercise 2.13. 🍷 *Show that the sublevel sets of a convex function are convex. Find a function with convex sublevel sets which is not convex.*

Exercise 2.14 (First order characterization of convex functions). 🍷🍷 *Let $f : \mathbb{R}^d \rightarrow \mathbb{R}$ be a differentiable function. Show that the following are equivalent:*

1. f is convex
2. f lies above its tangents: $\forall x, y \in \mathbb{R}^d, f(x) \geq f(y) + \langle \nabla f(y), x - y \rangle$
3. ∇f is monotone: $\forall x, y \in \mathbb{R}^d, \langle \nabla f(x) - \nabla f(y), x - y \rangle \geq 0$

Exercise 2.15 (Continuity). ☹☹☹ Show that a convex function is locally Lipschitz (hence continuous) on the interior of its domain.

Exercise 2.16 (Characterization of strongly convex functions in the Euclidean case). ☹ Show that f is μ -strongly convex with respect to the Euclidean norm if and only if $f - \frac{\mu}{2} \|\cdot\|^2$ is convex.

Exercise 2.17. ☹ Show that a strongly convex function admits exactly one minimizer.

2.2.2 Gradient

Exercise 2.18. ☹ Provide an example of setting where the gradient is not equal to the vector of partial derivatives.

Exercise 2.19. ☹ Show that the gradient of a function is orthogonal to the level lines of that function.

Exercise 2.20. ☹ Compute the gradients and Hessians of $x \mapsto \|x\|^2$, $x \mapsto \|x\|$, $x \mapsto a^\top x$. Is it true that the gradient of $x \mapsto \frac{1}{2}x^\top Ax$ is equal to Ax ?

Exercise 2.21. ☹ Let $A \in \mathbb{R}^{n \times d}$, $f : \mathbb{R}^n \rightarrow \mathbb{R}$ and $g : \mathbb{R}^d \rightarrow \mathbb{R}$ defined as $g(x) = f(Ax)$ for all $x \in \mathbb{R}^d$. Show that

$$\begin{aligned}\nabla g(x) &= A^* \nabla f(Ax) \ , \\ \nabla^2 g(x) &= A^* \nabla^2 f(Ax) A \ .\end{aligned}$$

Exercise 2.22. ☹☹ Compute the gradient of the logdet function, $M \mapsto \log \det(M)$.

3 Gradient descent on convex functions

In [Section 1](#) we proved that for gradient descent on ordinary least squares we can get very fast (so-called “linear”) convergence rates in terms of function values and iterates. Two numbers governed this convergence: global upper and lower bounds L and μ on the Hessian.

Our proof was very ad hoc, relying on the explicit expression of minimizer, the gradient and the Hessian. The questions we will try to answer in the sequel are:

- can we generalize the results of [Proposition 1.2](#) to other objective functions?
- under which conditions?
- what are the objective’s property that influence the convergence rate we obtain?

The first and weakest result we’ll prove in this class concerns convex Lipschitz functions.

Proposition 3.1 (Gradient descent on Lipschitz convex functions). *Let f be a convex, L -Lipschitz differentiable function admitting at least one minimizer x^* . For $t \in \mathbb{N}$, the iterates of gradient descent $x_{k+1} = x_k - \eta \nabla f(x_k)$ satisfy*

$$\min_{1 \leq k \leq t} f(x_k) - f(x^*) \leq \frac{L \|x_0 - x^*\|}{\sqrt{t}} , \quad (3.1)$$

and

$$f\left(\frac{1}{t} \sum_{k=1}^t x_k\right) - f(x^*) \leq \frac{L \|x_0 - x^*\|}{\sqrt{t}} , \quad (3.2)$$

if the stepsize η is taken as $\eta = \frac{\|x_0 - x^*\|}{L\sqrt{t}}$.

Let us make a few observations before the proof:

- This algorithm is a bit weird: the total number of iterations t must be known in advance to select the stepsize (in [Exercise 3.16](#) we get rid of this up to a slight worsening in the rate, using a decaying stepsize $\eta_k \propto 1/\sqrt{k}$).
- The objective values are not necessarily decreasing, it is not a *descent* algorithm.
- The more iterations we do, the smaller we need to take the stepsize.
- The stepsize depends on the unknown quantity $\|x_0 - x^*\|$. We can get rid of this dependency by bounding $\|x_0 - x^*\|$ with e.g. the diameter of the domain.
- The kind of rate is not as strong as in the OLS case: we know nothing about the last iterate x_t and only have results on the ergodic (averaged) iterates or on the best iterate.

Proof. Let x^* be a minimizer of f . For $k \in \mathbb{N}$, using convexity of f , definition of x_{k+1} and the parallelogram identity,

$$f(x_k) - f(x^*) \leq \langle \nabla f(x_k), x_k - x^* \rangle \quad (3.3)$$

$$\leq \frac{1}{\eta} \langle x_k - x_{k+1}, x_k - x^* \rangle \quad (3.4)$$

$$\leq \frac{1}{2\eta} (\|x_k - x_{k+1}\|^2 + \|x_k - x^*\|^2 - \|x_{k+1} - x^*\|^2) . \quad (3.5)$$

Summing, the telescopic series cancels out, and since f is L -Lipschitz we can bound $\|x_{k+1} - x_k\| = \eta \|\nabla f(x_k)\|$ by ηL , so:

$$\frac{1}{t} \sum_{k=1}^t f(x_k) - f(x^*) \leq \frac{1}{2t\eta} (t\eta^2 L^2 + \|x_0 - x^*\|^2 - \|x_{t+1} - x^*\|^2) \quad (3.6)$$

$$\leq \frac{\eta L^2}{2} + \frac{\|x_0 - x^*\|^2}{2t\eta} \quad (3.7)$$

Minimizing the RHS in η gives $\eta = \frac{R}{L\sqrt{t}}$ and the upper bound has value $\frac{RL}{\sqrt{t}}$. Jensen's inequality concludes. \square

The $1/\sqrt{k}$ rate is quite poor: to approach the optimal value within ε , one needs $\mathcal{O}(1/\varepsilon^2)$ iterations. With more assumptions on f , it can be improved to $\mathcal{O}(1/t)$, as we shall see in [Proposition 3.5](#).

Definition 3.2 (L -smoothness). *A differentiable function f is L -smooth if its gradient is L -Lipschitz: for all $x, y \in \text{dom } f$,*

$$\|\nabla f(x) - \nabla f(y)\| \leq L\|x - y\|. \quad (3.8)$$

A widely used property of L -smooth functions is that they satisfy the so-called Descent lemma.

Lemma 3.3 (Descent lemma). *Let f be a L -smooth function. Then for all $x, y \in \text{dom } f$,*

$$f(x) \leq f(y) + \langle \nabla f(y), x - y \rangle + \frac{L}{2} \|x - y\|^2. \quad (3.9)$$

Think of this inequality for a y fixed, and involving two functions of x , f and $\phi_y = f(y) + \langle \nabla f(y), \cdot - y \rangle + \frac{L}{2} \|\cdot - y\|^2$. The function ϕ_y is a convex, isotropic parabola. [Equation \(3.9\)](#) says that ϕ_y globally upper bounds f ; in addition, the two functions are equal and tangent at y .

Proof. Notice that:

$$f(x) - f(y) = \int_0^1 \frac{d}{dt} f(y + t(x - y)) dt = \int_0^1 \langle \nabla f(y + t(x - y)), x - y \rangle dt. \quad (3.10)$$

Subtract $\langle \nabla f(y), x - y \rangle$, use Cauchy-Schwarz and the hypothesis on f , conclude. \square

Lemma 3.4 (Cocoercivity of the gradient). *Let $f : \mathbb{R}^d \rightarrow \mathbb{R}$ be a L -smooth function.*

$$\forall x, y \in \mathbb{R}^d, \quad \langle \nabla f(x) - \nabla f(y), x - y \rangle \geq \frac{1}{L} \|\nabla f(x) - \nabla f(y)\|^2. \quad (3.11)$$

Proposition 3.5 (Convergence rate of gradient descent on L -smooth functions). *Let $f : \mathbb{R}^d \rightarrow \mathbb{R}$ be a convex L -smooth function with non empty set of minimizers. Then the iterates of gradient descent with stepsize $1/L$ converge at the rate:*

$$f(x_k) - f(x^*) \leq \frac{2L\|x_0 - x^*\|^2}{k}. \quad (3.12)$$

Proof. Let x^* be a minimizer of f . First, we show that the distance of x_k to x^* decreases:

$$\|x_{k+1} - x^*\|^2 = \|x_k - x^*\|^2 + 2\langle x_k - x^*, x_{k+1} - x_k \rangle + \|x_{k+1} - x_k\|^2 \quad (3.13)$$

$$= \|x_k - x^*\|^2 - \frac{2}{L}\langle \nabla f(x_k), x_k - x^* \rangle + \frac{1}{L^2}\|\nabla f(x_k)\|^2. \quad (3.14)$$

Lemma 3.4, combined with $\nabla f(x^*) = 0$, yields:

$$-\frac{2}{L}\langle \nabla f(x_k), x_k - x^* \rangle + \frac{1}{L^2}\|\nabla f(x_k)\|^2 \leq -\frac{1}{L^2}\|\nabla f(x_k)\|^2 \leq 0, \quad (3.15)$$

which concludes: $\|x_{k+1} - x^*\| \leq \|x_k - x^*\|$.

We then use convexity of f :

$$f(x_k) - f(x^*) \leq \langle \nabla f(x_k), x_k - x^* \rangle \leq \|\nabla f(x_k)\| \|x_k - x^*\| \leq \|\nabla f(x_k)\| \|x_0 - x^*\|. \quad (3.16)$$

Finally the descent Lemma (**Lemma 3.3**) quantifies the decrease in objective at each iteration:

$$f(x_{k+1}) - f(x_k) \leq -\frac{1}{2L}\|\nabla f(x_k)\|^2. \quad (3.17)$$

Introducing $\delta_k = f(x_k) - f(x^*)$, we have:

$$\delta_{k+1} - \delta_k \leq -\frac{1}{2L}\|\nabla f(x_k)\|^2 \quad (3.18)$$

$$\leq -\frac{1}{2L} \frac{\delta_k^2}{\|x_0 - x^*\|^2}. \quad (3.19)$$

The above manipulations are quite standard in optimization, but the following trick is a bit surprising the first time you see it: dividing by $\delta_k \delta_{k+1}$,

$$\frac{1}{\delta_k} - \frac{1}{\delta_{k+1}} \leq -\frac{1}{2L\|x_0 - x^*\|^2} \frac{\delta_k}{\delta_{k+1}} \quad (3.20)$$

$$\leq -\frac{1}{2L\|x_0 - x^*\|^2}, \quad (3.21)$$

since (δ_k) decreases (by (3.18)). Summing and telescoping concludes. \square

If the function is even more regular, the rate of gradient descent can further be improved.

Proposition 3.6. *Let $f : \mathbb{R}^d \rightarrow \mathbb{R}$ be L -smooth and μ -strongly convex. Then it admits a unique minimizer x^* (**Exercise 2.17**) and the iterates of gradient descent with stepsize $1/L$ converge linearly:*

$$f(x_k) - f(x^*) \leq \left(1 - \frac{\mu}{L}\right)^k (f(x_0) - f(x^*)). \quad (3.22)$$

Proof. Let x^* be the minimizer of f . Being μ -strongly convex, f satisfies the Polyak-Łojasiewicz inequality (**Exercise 3.11**):

$$\forall x \in \mathbb{R}^d, f(x) - f(x^*) \leq \frac{1}{2\mu}\|\nabla f(x)\|^2. \quad (3.23)$$

Combined with the descent lemma,

$$f(x_{k+1}) - f(x_k) \leq -\frac{1}{2L} \|\nabla f(x_k)\|^2 \quad (3.24)$$

$$f(x_{k+1}) - f(x^*) + f(x^*) - f(x_k) \leq -\frac{\mu}{L} (f(x_k) - f(x^*)) \quad (3.25)$$

hence $f(x_{k+1}) - f(x^*) \leq (1 - \frac{\mu}{L})(f(x_k) - f(x^*))$. □

3.1 Exercises

3.1.1 Convexity inequalities

Exercise 3.7. ☹☹ Let f be a convex and differentiable function. Let $L > 0$. Show that the following properties are equivalent:

1. $\forall x, y, \|\nabla f(x) - \nabla f(y)\| \leq L\|x - y\|$
2. $\forall x, y, f(x) \leq f(y) + \langle \nabla f(y), x - y \rangle + \frac{L}{2}\|x - y\|^2$
3. $\forall x, y, \frac{1}{L}\|\nabla f(x) - \nabla f(y)\|^2 \leq \langle x - y, \nabla f(x) - \nabla f(y) \rangle$

Exercise 3.8. ☹ Let f be a twice differentiable L -smooth function. Show that for all $x \in \mathbb{R}^d$, $\nabla^2 f(x) \preceq L \text{Id}$.

Exercise 3.9. ☹ Let f be a differentiable function. Show that the following properties are equivalent:

1. f is μ -strongly convex
2. $\forall x, y, f(x) \geq f(y) + \langle \nabla f(y), x - y \rangle + \frac{\mu}{2}\|x - y\|^2$
3. $\forall x, y, \mu\|x - y\|^2 \leq \langle x - y, \nabla f(x) - \nabla f(y) \rangle$

Exercise 3.10. ☹ Let f be a twice differentiable μ -strongly convex function. Show that for all $x \in \mathbb{R}^d$, $\mu \text{Id} \preceq \nabla^2 f(x)$.

Exercise 3.11 (Polyak-Łojasiewicz inequality). ☹☹ Let f be a μ -strongly-convex and differentiable function. Let $x^* = \text{argmin} f(x)$. Show that f satisfies the Polyak-Łojasiewicz inequality:

$$\mu(f(x) - f(x^*)) \leq \frac{1}{2}\|\nabla f(x)\|^2 .$$

Provide an example of function which is not strongly convex, but satisfies the inequality.

Exercise 3.12. ☹☹ Let f be a L -smooth μ -strongly convex function. Show that for any x, y ,

$$\frac{\mu L}{\mu + L}\|x - y\|^2 + \frac{1}{L + \mu}\|\nabla f(x) - \nabla f(y)\|^2 \leq \langle x - y, \nabla f(x) - \nabla f(y) \rangle .$$

Exercise 3.13 (3 point descent lemma). ☹ Let f be convex and L -smooth. Show that for any triplet (x, y, z) ,

$$f(x) \leq f(y) + \langle \nabla f(z), x - y \rangle + \frac{L}{2}\|x - z\| .$$

3.1.2 Gradient descent

Exercise 3.14 (?? in discrete time). ☹☹ Let $f : \mathbb{R}^d \rightarrow \mathbb{R}$ be a convex, twice differentiable L -smooth function.

Show that the iterates of gradient descent on f with step size $0 < \alpha < 2/L$ have decreasing gradient norm.

Can you show it if f is not twice differentiable?

Is it still true when f is not convex?

Exercise 3.15. ☹ Provide a finite dimensional example of convex L -smooth function f such that gradient descent with stepsize $< 2/L$ diverges.

Exercise 3.16 (Gradient descent on Lipschitz function without knowing the horizon). ☹☹ For gradient descent on a Lipschitz differentiable objective, the classical proof assumes that the total number of iterations t is known in advance, to set the fixed stepsize $\eta \propto 1/\sqrt{t}$. Show that using a decreasing stepsize $\eta_k \propto 1/\sqrt{k}$ leads to a rate of order $\log k/\sqrt{k}$ on the ergodic or best iterate.

4 Non smooth convex optimization

For a differentiable objective function $f : \mathbb{R}^d \rightarrow \mathbb{R}$, we have seen:

1. convexity: f is above its tangents
2. Lipschitzness
3. L -smoothness, implying the descent lemma (Lemma 3.3), meaning that at any point $y \in \mathbb{R}^d$ there exists an isotropic parabola of curvature L that globally **upper** bounds f and is tangent to f at y .
4. strong convexity: at any point $y \in \mathbb{R}^d$ there exists an isotropic parabola of curvature μ that globally **lower** bounds f and is tangent to f at y .

Depending on the properties satisfied by f , we have:

- 1 gives the necessary and sufficient condition for x to be a minimizer: $\nabla f(x^*) = 0$.
- 1 + 2 gives a convergence rate of $1/\sqrt{k}$ under stepsize depending on the total number of iterations. The rate is in objective, for the best iterate or the averaged iterates.
- 1 + 3 improves it to $1/k$ with stepsize $1/L$ not depending on the number of iterations. The rate is in objective, on the last iterate.
- 3 + 4 sandwiches the functions between two isotropic parabolas of curvature μ and L and gives the excellent linear rate. The rate is on the distance from last iterate to the unique minimizer.

Now we want to consider non-differentiable functions f . Why do we need to handle these?

4.1 Constrained optimization

Optimizers sometimes wants the minimizer of their cost functions to belong to some set $\mathcal{C} \subset \mathbb{R}^d$.

For example, suppose that the variable x corresponds to an image, stored as a 2D array of pixels – an element of $\mathbb{R}^{d \times d}$. Then the values of the pixels are not just any real numbers: they should be in $[0, 1]$ if they correspond to gray levels for example.

If the optimization variable $x \in \mathbb{R}^d$ represents proportions (say, of various types of cells in a sane/cancerous tissue), it should have non-negative entries that sum to 1: $x \in \Delta_d$, where Δ_d is the d -dimensional *simplex*

$$\Delta_d = \{x \in \mathbb{R}^d : \forall i \in [d], x_i \geq 0, \sum_{i=1}^d x_i = 1\} . \quad (4.1)$$

Finally, suppose the practitioner is looking for the approximate solution to a linear system (using least squares), but wants a solution x^* that uses as few variables as possible. There are many ways to do this, but one popular (because convex) way to do so is to constrain the ℓ_1 norm of x :

$$\min_{x \in \mathbb{R}^d} \frac{1}{2} \|Ax - b\|^2 \quad \text{subject to} \quad \|x\|_1 \leq \tau .$$

All these examples lead us to consider constrained problems:

$$\min_{x \in \mathcal{C}} f(x) . \quad (4.2)$$

Because of the constraint, even if f is still smooth, the tools we have used so far no longer apply. Consider the basic one-dimensional problem:

$$\min_{x \in [0,1]} x . \quad (4.3)$$

This problem is friendly: the objective is convex, the constraint set $[0,1]$ is convex and compact. The minimizer is 0, yet $\nabla f(0) = 1 \neq 0$ and so for constrained convex optimization, the global characterization of minimizers ([Proposition 2.6](#)) does not hold.

Fortunately, we can replace some concepts seen so far by generalizations that are very well adapted. First, let us introduce a tool that allows to remove the constraints.

4.2 Extended value functions

In optimization, it is often convenient to work with functions that take values in $\overline{\mathbb{R}} \triangleq \mathbb{R} \cup \{+\infty\}$. The prototypical example is the *indicator* function of a set.

Definition 4.1 (Indicator function). *In convex analysis, the indicator function $\iota_{\mathcal{C}}$ of a subset \mathcal{C} of \mathbb{R}^d is:*

$$\begin{aligned} \iota_{\mathcal{C}} : \mathbb{R}^d &\rightarrow \overline{\mathbb{R}} \\ x &\mapsto \begin{cases} 0 , & \text{if } x \in \mathcal{C} , \\ +\infty , & \text{otherwise} . \end{cases} \end{aligned} \quad (4.4)$$

Note that this is different from the other indicator function that takes value 1 on the set and 0 outside⁴. The indicator function conveniently allows transforming all constrained minimization problems into unconstrained ones of course at the price of working with extended value functions:

$$\operatorname{argmin}_{x \in \mathcal{C}} f(x) = \operatorname{argmin}_{x \in \mathbb{R}^d} f(x) + \iota_{\mathcal{C}}(x) . \quad (4.5)$$

Definition 4.2. *The domain of a function $f : \mathbb{R}^d \rightarrow \overline{\mathbb{R}}$ is:*

$$\operatorname{dom} f \triangleq \{x \in \mathbb{R}^d, f(x) < +\infty\} . \quad (4.6)$$

The second tool to overcome non-differentiability is a substitute for the gradient.

4.3 Subdifferential of convex functions

Definition 4.3. *The subdifferential of f at x is the set of slopes of all affine minorants of f that are exact at x :*

$$\partial f(x) = \{u \in \mathbb{R}^d : \forall y \in \mathbb{R}^d, f(y) \geq f(x) + \langle u, y - x \rangle\} . \quad (4.7)$$

An element of the subdifferential is called a subgradient.

⁴this poor function is rarely convex

Note that contrary to the gradient, the subdifferential is a set-valued mapping: the subdifferential at one point may contain more than one subgradient. It may also be empty at some point ([Exercise 4.16](#)), but it's not on the interior of the domain if f is convex ([Exercise 4.22](#)).

It is a generalization of the gradient in the following sense.

Proposition 4.4. *Let f be a convex differentiable function. Then the subdifferential only contains the gradient:*

$$\partial f(x) = \{\nabla f(x)\} . \quad (4.8)$$

Note: the converse is true (if the subdifferential reduces to a point, f is differentiable at this point).

The subdifferential allows to elegantly characterize the minimizers of *all* convex functions, even the nondifferentiable ones.

Proposition 4.5 (Fermat's rule). *Let f be convex. Then for all x ,*

$$x \in \operatorname{argmin} f \Leftrightarrow 0 \in \partial f(x) . \quad (4.9)$$

The proof is 1 line and left to the reader.

Example 4.6. *The subdifferential of the absolute value is:*

$$\partial |\cdot|(x) = \begin{cases} \{x/|x|\}, & \text{if } x \neq 0 , \\ [-1, 1], & \text{otherwise} . \end{cases} \quad (4.10)$$

What is the subdifferential of $\iota_{[0,1]}$?

Proposition 4.7 (Subdifferential of sum). *The subdifferential of the sum contains the sum of subdifferentials:*

$$\partial(f + g) \supset \partial f + \partial g , \quad (4.11)$$

(to be understood as the Minkowski sum). Equality does not hold in general, although the counter examples are pathological ([Exercise 4.23](#)).

Remark 4.8 (Subgradient descent). *If you look again with nonsmooth eyes at [Proposition 3.1](#) regarding gradient descent on Lipschitz convex functions, we actually never used the fact that f was differentiable: we only started with $f(x_k) - f(x^*) \leq \langle \nabla f(x_k), x_k - x^* \rangle$. Since a subgradient of f at x_k , by definition, also satisfies this inequality, it means we can replace $\nabla f(x_k)$ by any $g_k \in \partial f(x_k)$ and the proof still holds.*

This algorithm is called subgradient descent.

The last tool we need for nonsmooth optimization is the proximal operator.

4.4 The proximal operator

So far we have brushed away the existence of minimizers, and only assumed they exist. In classical analysis, the celebrated Weierstrass theorem says that a continuous function is bounded and reaches its minimum on any compact. We will provide a slight relaxation of this theorem that handles a slightly weaker property than continuity.

Definition 4.9 (Lower semi-continuity). *The function $f : \mathbb{R}^d \rightarrow \overline{\mathbb{R}}$ is lower semicontinuous at $x \in \text{dom } f$ if for any sequence (x_k) converging to x ,*

$$f(x) \leq \liminf_{k \rightarrow \infty} f(x_k) . \quad (4.12)$$

Any continuous function is clearly lower semicontinuous.

Remark 4.10. *Lower semicontinuous functions are also referred to as closed functions. This is because a function is l.s.c. if and only if its epigraph*

$$\text{epi } f = \{(x, t) : f(x) \leq t\} \subset \mathbb{R}^d \times \mathbb{R} \quad (4.13)$$

is closed.

Proposition 4.11 (Existence of minimizers, the direct method of calculus of variations). *Let $f : \mathbb{R}^d \rightarrow \overline{\mathbb{R}}$ be a coercive lower semicontinuous function, with non empty domain. Then f admits at least one minimizer.*

Note: this result is only valid in finite dimension.

Proof. Let x_0 such that $f(x_0) < +\infty$. Since f is coercive, $f(x) \rightarrow_{\|x\| \rightarrow \infty} +\infty$. Let M such that $\|x\| \geq M \implies f(x) \geq f(x_0)$. Let \mathcal{B} be the ball of center 0, radius M . Since we are in finite dimension, \mathcal{B} is compact.

Let $f^* = \inf_{x \in \mathcal{B}} f(x)$. In general, we could have $f^* = -\infty$ (think $d = 1$, $M = 1$, $f(x) = 1/x$ except $f(0) = 0$). But we'll show that since f is l.s.c., it can't be the case.

Assume $f^* = -\infty$. We can construct a sequence (x_k) in \mathcal{B} such that $\forall k \in \mathbb{N}$, $f(x_k) \leq -k$. But \mathcal{B} is compact: we can extract a converging subsequence, that we call y_k , converging to $\tilde{x} \in \mathcal{B}$.

By lower semicontinuity of f , $f(\tilde{x}) \leq \liminf_{k \rightarrow \infty} f(y_k) = -\infty$ which is not possible since f does not take value $-\infty$.

Hence f^* is finite. Now we do the exact same reasoning to construct (x_k) such that $f(x_k) \leq f^* + 1/k$. Take a converging subsequence, use lower semicontinuity to show that the limit point \tilde{x} satisfies $f(\tilde{x}) \leq f^*$ and thus $f(\tilde{x}) = f^*$. \square

Definition 4.12. *The function $f : \mathbb{R}^d \rightarrow \overline{\mathbb{R}}$ is said to be (or more rigorously belong to) $\Gamma_0(\mathbb{R}^d)$ if it is:*

- *convex*
- *proper (its domain is not empty)*
- *lower semicontinuous*

When is $\iota_{\mathcal{C}} \Gamma_0$?

Definition 4.13 (Proximal operator). *Let $f \in \Gamma_0(\mathbb{R}^d)$. The proximal operator of f evaluated at x is*

$$\text{prox}_f(x) = \underset{y \in \mathbb{R}^d}{\operatorname{argmin}} f(y) + \frac{1}{2} \|x - y\|^2 . \quad (4.14)$$

Why is it well-defined (why does the infimum exist? why is it attained? why does the argmin contain exactly one point)?

The following is a useful characterization of proxs.

Proposition 4.14 (Characterization of proximal operators). *Let $f \in \Gamma_0(\mathbb{R}^d)$. Then $p = \text{prox}_f(x)$ if and only if $x - p \in \partial f(p)$.*

This justifies the frequent formulation $\text{prox}_f = (\text{Id} + \partial f)^{-1}$ (a.k.a. the prox is the resolvent of the subdifferential).

Proof. Apply Fermat's rule. □

Proximal operators may seem new, but without realizing it you know and have used some particular instances: if \mathcal{C} is non empty, closed and convex, $\text{prox}_{\iota_{\mathcal{C}}}$ is the (well-defined) projection onto \mathcal{C} . For this and other reasons, proxs can be thought of as generalization of projections.

4.5 The proximal point algorithm

Let's jump ahead to the continuous time version of gradient descent that we'll see later, the gradient flow. The gradient flow is an Ordinary Differential Equation (ODE):

$$\begin{aligned} x(0) &= x_0, \\ \dot{x}(t) &= -\nabla f(x(t)) . \end{aligned} \quad (4.15)$$

Suppose one wants to numerically approximate the solution to this equation. We pick a time discretization step $\eta > 0$, define time instants $t_k = k\eta$ for $k \in \mathbb{N}$. We will construct a sequence x_k that should approximate the continuous version $x(t_k)$. We approximate the derivative $\dot{x}(t_k)$ by finite differences $\frac{x_{k+1} - x_k}{\eta}$.

Thus a discrete approximation of (4.15) is:

$$x_{k+1} - x_k = -\eta \nabla f(x_k) ; \quad (4.16)$$

this is a *forward* Euler discretization scheme, that gives the iterations of gradient descent.

On the other hand, one could chose a *backward* Euler scheme, also called *implicit* (because it does not give an explicit value of x_k):

$$x_{k+1} - x_k = -\eta \nabla f(x_{k+1}) . \quad (4.17)$$

Rewriting this as $(\text{Id} + \eta \nabla f)(x_{k+1}) = x_k$, it becomes visible that this scheme is the proximal point algorithm!

By Fermat's rule and the characterization of proxs ([Proposition 4.14](#)), x^* is a minimizer of f if and only if it is a fixed point of prox_f : $x^* = \text{prox}_f(x^*)$.

Proposition 4.15 (Firm non expansivity of proximal operators). *Proximal operators are firmly non expansive:*

$$\|x - y\|^2 \leq \langle \text{prox}_f(x) - \text{prox}_f(y), x - y \rangle . \quad (4.18)$$

By Cauchy-Schwarz inequality, this immediately implies that proxs are non-expansive:

$$\| \text{prox}_f(x) - \text{prox}_f(y) \| \leq \|x - y\| . \quad (4.19)$$

Proof. Let $p_x = \text{prox}_f(x)$, $p_y = \text{prox}_f(y)$. The prox characterization for p_x is $x - p_x \in \partial f(p_x)$, so by definition of the subdifferential:

$$f(p_y) \geq f(p_x) + \langle x - p_x, p_y - p_x \rangle . \quad (4.20)$$

Do the same for p_y , sum the two inequations to cancel out the f 's, reorder. \square

4.6 The proximal gradient descent algorithm

A forward backward splitting. View as a MM approach. Convergence rate.

4.7 Exercises

4.7.1 Subdifferential

Exercise 4.16. ☹ Provide an example of convex function which has an empty subdifferential at some point of its domain.

Exercise 4.17. ☹ Let $f : \mathbb{R}^d \rightarrow \mathbb{R}$ be separable: $f(x) = \sum_1^d f_i(x_i)$ where the f_i 's are functions of the real variable.

Show that $\partial f(x) = \partial f_1(x_1) \times \dots \times \partial f_d(x_d)$.

Compute the subdifferential of the ℓ_1 -norm.

Exercise 4.18. ☹ Show that the subdifferential of a function at a point writes as an intersection of half-spaces. Show that it is convex and closed.

Exercise 4.19. ☹ Show that the subdifferential of ι_C at $x \in C$ is equal to the normal cone of C at x , that is:

$$\mathcal{N}_C(x) = \{u : \forall z \in C, \langle u, z - x \rangle \leq 0\}$$

Exercise 4.20. ☹☹ Compute the subdifferential of the Euclidean norm at any point in \mathbb{R}^d .

Exercise 4.21 (Exercise 4.25 revisited). ☹☹ Show that the first order optimality condition for differentiable convex constrained optimization rewrites:

$$x^* \in \underset{C}{\operatorname{argmin}} f(x) \Leftrightarrow -\nabla f(x^*) \in \mathcal{N}_C(x^*) .$$

Exercise 4.22 (Non emptiness of subdifferential). ☹☹☹ Let $f : \mathbb{R}^d \rightarrow \overline{\mathbb{R}}$ be convex. Show that if $x \in \operatorname{int}(\operatorname{dom} f)$, $\partial f(x)$ is non empty and compact.

Show that if x lies on the boundary of $\operatorname{dom} f$, $\partial f(x)$ is either empty or unbounded.

Exercise 4.23. ☹ This exercise is in dimension 2. Let \mathcal{C}_1 and \mathcal{C}_2 be two closed balls of strictly positive radius, that share a single point. Show that $\partial(\iota_{\mathcal{C}_1} + \iota_{\mathcal{C}_2})$ is a strict superset of $\partial\iota_{\mathcal{C}_1} + \iota_{\mathcal{C}_2}$ at this point.

4.7.2 Constrained optimization

Exercise 4.24. ☹ Show that the indicator function ι_C is convex (resp. lower semicontinuous, resp. proper) if C is convex (resp. closer, resp. nonempty).

Exercise 4.25 (Global optimality condition for constrained convex optimisation). ☹ Let $f : \mathbb{R}^d \rightarrow \mathbb{R}$ be a convex differentiable function, let C be a convex subset of \mathbb{R}^d . Show that $x^* \in \operatorname{argmin}_{x \in C} f(x)$ if and only if

$$\forall x \in C, \langle \nabla f(x^*), x - x^* \rangle \geq 0 .$$

Exercise 4.26. ☹☹ Let C be a non-empty closed convex set. Show that for all $x \in C, y \in \mathbb{R}^d$,

$$\|x - y\|^2 \geq \|x - \Pi_C(y)\|^2 + \|y - \Pi_C(y)\|^2 .$$

In particular this shows that $\|y - x\| \geq \|x - \Pi_C(y)\|$, which says that projection can only get you closer to optimum (taking $x = x^*$, if your current iterate is y).

5 Duality

One motivation (not the single reason for this chapter to exist!): when do we stop iterative algorithms? Usually we'd settle for $f(x_k) - f(x^*)$ less than some tolerance, but $f(x^*)$ is unknown.

5.1 The Fenchel transform

A fundamental tool in convex analysis is the Fenchel transform (sometimes called the Fenchel-Legendre or Legendre transform), that is, in many ways, akin to the Fourier transform in signal processing.

Definition 5.1 (Fenchel conjugate). *Let $f : \mathbb{R}^d \rightarrow \overline{\mathbb{R}}$. The Fenchel conjugate (or transform) of f is:*

$$f^*(u) = \sup_{x \in \mathbb{R}^d} \langle u, x \rangle - f(x) . \quad (5.1)$$

It is always convex, even if f is not ([Exercise 2.8](#)). To get familiar with it, compute the Fenchel transforms of $\|\cdot\|^2/2a$ ($a > 0$), $\|\cdot\|$ and $x \mapsto \sup_{y \in [-1,1]} x$.

Geometrical interpretation of the Fenchel transform $f^*(u) \leq \alpha$ is equivalent to

$$\forall x \in \mathbb{R}^d, \quad \langle x, u \rangle - \alpha \leq f(x) , \quad (5.2)$$

meaning that $-\alpha$ is the intercept of a global affine minorant of f of slope u . $-f^*(u)$ is thus the biggest possible intercept of such minorants – beware, there may not exist such minorants ($f^*(u) = +\infty$).

Proposition 5.2 (Involution on Γ_0). *If $f \in \Gamma_0(\mathbb{R}^d)$, $f^{**} = f$.*

Proposition 5.3 (Fenchel-Young inequality). *Let $f : \mathbb{R}^d \rightarrow \overline{\mathbb{R}}$. Then for all $x, u \in \mathbb{R}^d$,*

$$f(x) + f^*(u) \geq \langle x, u \rangle . \quad (5.3)$$

Equality holds if and only if $u \in \partial f(x)$.

Proof. The first part is trivial. Then,

$$\begin{aligned} f^*(u) = \langle x, u \rangle - f(x) &\Leftrightarrow x \in \operatorname{argmin} f - \langle \cdot, u \rangle \\ &\Leftrightarrow 0 \in \partial(f - \langle \cdot, u \rangle)(x) \\ &\Leftrightarrow u \in \partial f(x) . \end{aligned}$$

□

Proposition 5.4. *If $f \in \Gamma_0(\mathbb{R}^d)$, $\partial f^* = (\partial f)^{-1}$, in the sense that:*

$$u \in \partial f(x) \Leftrightarrow x \in \partial f^*(u) . \quad (5.4)$$

Proof. Use then Fenchel-Young equality for f and then use $f^{**} = f$. □

Proposition 5.5 (Moreau decomposition formula). *Let $f \in \Gamma_0(\mathbb{R}^d)$. Then*

$$\text{prox}_{\gamma f} + \gamma \text{prox}_{\gamma^{-1} f^*}(\cdot/\gamma) = \text{Id}$$

Proof.

$$p = \text{prox}_{\gamma f}(x) \Leftrightarrow \frac{x - p}{\gamma} \in \partial f(p) \quad (5.5)$$

$$\Leftrightarrow p \in \partial f^*\left(\frac{x - p}{\gamma}\right) \quad (5.6)$$

$$\Leftrightarrow \frac{x}{\gamma} - \frac{x - p}{\gamma} \in \partial \frac{1}{\gamma} f^*\left(\frac{x - p}{\gamma}\right) \quad (5.7)$$

$$\Leftrightarrow \frac{x - p}{\gamma} = \text{prox}_{f^*/\gamma}\left(\frac{x}{\gamma}\right) . \quad (5.8)$$

□

Proposition 5.6. *f is μ -strongly convex iff f^* is $\frac{1}{\mu}$ -smooth.*

5.2 Fenchel-Rockafellar duality

Many problems in Machine Learning have the following form:

$$\min_{x \in \mathbb{R}^d} f(Ax) + g(x) = P(x) , \quad (5.9)$$

with $A \in \mathbb{R}^{n \times d}$, $f \in \Gamma_0(\mathbb{R}^n)$ and $g \in \Gamma_0(\mathbb{R}^d)$.

Definition 5.7. *The Fenchel-Rockafellar dual problem of [Problem \(9\)](#) is:*

$$\min_{y \in \mathbb{R}^n} f^*(y) + g^*(-A^\top y) = D(y) . \quad (5.10)$$

It is often alternatively written as a concave maximization problem

$$\max_{y \in \mathbb{R}^n} -f^*(y) - g^*(-A^\top y) . \quad (5.11)$$

Where does this problem come from? It can be obtained by Lagrangian duality:

$$\inf_{x \in \mathbb{R}^d} f(Ax) + g(x) = \inf_{x \in \mathbb{R}^d, z \in \mathbb{R}^n} f(z) + g(x) \quad \text{s.t. } Ax = z \quad (5.12)$$

$$= \inf_{x \in \mathbb{R}^d, z \in \mathbb{R}^n} f(z) + g(x) + \iota_{\{z\}}(Ax) \quad (5.13)$$

$$= \inf_{x \in \mathbb{R}^d, z \in \mathbb{R}^n} f(z) + g(x) + \sup_{y \in \mathbb{R}^n} \langle y, Ax - z \rangle \quad (5.14)$$

$$= \inf_{x \in \mathbb{R}^d, z \in \mathbb{R}^n} \sup_{y \in \mathbb{R}^n} f(z) + g(x) + \langle y, Ax - z \rangle \quad (5.15)$$

Let's swap inf and sup. By doing so, we lose the equivalence ($\sup \inf \leq \inf \sup$, without equality in general). We have:

$$\sup_{y \in \mathbb{R}^n} \inf_{x \in \mathbb{R}^d, z \in \mathbb{R}^n} f(z) + g(x) + \langle y, Ax - z \rangle \quad (5.16)$$

$$= \sup_{y \in \mathbb{R}^n} \inf_{x \in \mathbb{R}^d} g(x) + \langle y, Ax \rangle + \inf_{z \in \mathbb{R}^n} f(z) - \langle y, z \rangle \quad (5.17)$$

$$= \sup_{y \in \mathbb{R}^n} - \sup_{x \in \mathbb{R}^d} -g(x) + \langle -A^\top y, x \rangle - \sup_{z \in \mathbb{R}^n} -f(z) + \langle y, z \rangle \quad (5.18)$$

$$= \sup_{y \in \mathbb{R}^n} -g^*(-A^\top y) - f^*(y) , \quad (5.19)$$

that is our dual problem.

Proposition 5.8 (Weak duality). *For all $x, y \in \mathbb{R}^d \times \mathbb{R}^n$, $P(x) + D(y) \geq 0$.*

Proof. Write the sum, group f with f^* and g with g^* , then use the Fenchel-Young inequality. \square

Proposition 5.9. *Let $x^*, y^* \in \mathbb{R}^d \times \mathbb{R}^n$. The following conditions are equivalent:*

1. $x^* \in \operatorname{argmin} P$, $y^* \in \operatorname{argmin} D(y)$, $\inf P + \inf D = 0$.
2. $P(x^*) + D(y^*) = 0$
3. $f(Ax^*) + f^*(y^*) = \langle Ax^*, y^* \rangle$ and $g(x^*) + g^*(-A^\top y^*) = -\langle x^*, A^\top y^* \rangle$,
4. $-A^\top y^* \in \partial g(x^*)$ and $y^* \in \partial f(Ax^*)$
5. $x^* \in \partial g^*(-A^\top y^*)$ and $Ax^* \in \partial f^*(y^*)$

Proof. 1 same as 2 is OK.

2 same as 3: Do the same as in the proof of [Proposition 5.8](#), and use that the sum of two positive terms is 0, if and only if both terms are 0.

3 same as 4: equality in Fenchel-Young case

4 same as 5: [Proposition 5.4](#) \square

Definition 5.10 (Saddle point of the Lagrangian). *The pair $x^*, y^* \in \mathbb{R}^d \times \mathbb{R}^n$ is said to be a saddle-point of the Lagrangian if:*

$$\forall x, y \in \mathbb{R}^d \times \mathbb{R}^n, \mathcal{L}(x^*, y^*) \leq \mathcal{L}(x, y^*) \leq \mathcal{L}(x, y^*) . \quad (5.20)$$

Proposition 5.11. *If $x^*, y^* \in \mathbb{R}^d \times \mathbb{R}^n$ is said to be a saddle-point of the Lagrangian, then x^* solves the primal and y^* solves the dual.*

Proof. Write the two inequalities stating that the pair is a saddle point, make subgradients appear, use [Proposition 5.9](#). \square

Definition 5.12 (Qualification condition). *The following inclusion is called a qualification condition:*

$$0 \in \operatorname{int}(\operatorname{dom} f - A \operatorname{dom} g) . \quad (5.21)$$

Where does it come from? It can be seen as a stronger version of “the primal objective P is proper”:

$$\begin{aligned}
\text{dom } P \neq \emptyset &\Leftrightarrow \exists x \in \mathbb{R}^d, f(Ax) + g(x) < +\infty \\
&\Leftrightarrow \exists x \in \mathbb{R}^d, f(Ax) < +\infty, g(x) < +\infty \\
&\Leftrightarrow \exists x \in \mathbb{R}^d, y \in \mathbb{R}^n, y = Ax, f(y) < +\infty, g(x) < +\infty \\
&\Leftrightarrow \exists x \in \mathbb{R}^d, y \in \mathbb{R}^n, y = Ax, y \in \text{dom } f, x \in \text{dom } g \\
&\Leftrightarrow \exists x \in \text{dom } g, y \in \text{dom } f, 0 = y - Ax \\
&\Leftrightarrow 0 \in \text{dom } f - A \text{dom } g .
\end{aligned} \tag{5.22}$$

Proposition 5.13 (Strong duality). *Suppose that QC holds. Then there exist a saddle point of the Lagrangian, and $\inf P + \inf D = 0$ (the duality gap is 0 at optimum).*

Proof. The proof is not easy. □

Proposition 5.14.

5.3 Primal dual algorithms

In this section we introduce more algorithms to deal with composite problems, based on duality/splitting.

Chambolle-Pock algorithm The Chambolle-Pock algorithm solves problems of the form:

$$\min_{x \in \mathbb{R}^d} f(Ax) + g(x) , \tag{5.23}$$

where $f \in \Gamma_0(\mathbb{R}^n)$, $g \in \Gamma_0(\mathbb{R}^d)$, but both need not be smooth. It requires access to their proximal operators. Let σ and τ be positive stepsizes such that $\sigma\tau\|A\|_2^2 < 1$. Let $\eta \in]0, 1]$. The Chambolle-Pock iterations write:

$$\begin{aligned}
\bar{y}_k &= y_k + \eta(y_k - y_{k-1}) \\
x_{k+1} &= \text{prox}_{\tau g}(x_k - \tau A^\top \bar{y}_k) \\
y_{k+1} &= \text{prox}_{\sigma f^*}(y_k + \sigma Ax_k)
\end{aligned} \tag{5.24}$$

Note: The interpolation parameter η is most often taken equal to 1. The interpolation step can be performed on x instead of y , which yields a different version of the algorithm. The extension proposed by Vu and Condat handles an additional smooth term h in the objective.

ADMM The Alternating Direction Method of Multipliers (ADMM method) is designed to solve problems of the form:

$$\min_{x \in \mathbb{R}^d, z \in \mathbb{R}^n} f(z) + g(x) \quad \text{s.t.} \quad Ax + Bz = c . \tag{5.25}$$

For simplicity and to match the setup of Chambolle-Pock we take $B = -\text{Id}$ and $c = 0$.

The augmented Lagrangian of parameter ρ is:

$$\mathcal{L}^\rho(x, y, z) = f(z) + g(x) + \langle y, Ax - z \rangle + \frac{\rho}{2} \|Ax - z\|^2 \quad (5.26)$$

The algorithm reads:

$$\begin{aligned} x_{k+1} &\in \operatorname{argmin}_x \mathcal{L}^\rho(x, y_k, z_k) \\ y_{k+1} &\in \operatorname{argmin}_y \mathcal{L}^\rho(x_{k+1}, y, z_k) \\ y_{k+1} &= y_k + r(Ax_{k+1} - z_{k+1}) \end{aligned} \quad (5.27)$$

5.4 Exercises

Exercise 5.15. ☞ Show that the Fenchel transform is order reversing: if $f \leq g$, $gg \leq f^*$.

Exercise 5.16. ☞☞ Compute the Fenchel transform of the ℓ_p -norm for $p \in [1, +\infty]$.

Exercise 5.17. ☞☞ Show that $x \mapsto \frac{1}{2}\|x\|^2$ is a fixed point of the Fenchel transform. Show that it is the only fixed point of the Fenchel transform.

Exercise 5.18. ☞ Let $f : \mathbb{R}^d \rightarrow \mathbb{R}$ be separable: $f(x) = \sum_1^d f_i(x_i)$ where the f_i 's are functions of the real variable.

Show that $f^*(u) = (f_i^*(u_i))_{i \in [d]}$.

Exercise 5.19 (“Lipschitzing trick”). ☞☞ Let $f \in \Gamma_0(\mathbb{R}^d)$.

Show that f is M -Lipschitz if and only if the domain of f^* is included in the Euclidean ball of center 0 and radius M .

Exercise 5.20 (Convolution smoothing). ☞ Let $f : \mathbb{R}^d \rightarrow \mathbb{R}$ be convex.

Show that $f \square \frac{L}{2}\|\cdot\|^2$ is L -smooth.

6 Second order optimization: Newton method

Let f be a twice differentiable function. We have seen that gradient descent for stepsize $1/L$ on a L smooth function can be interpreted as a Majorization-Minimization approach, using the descent lemma to upper bound f globally by an isotropic parabola (majorization step), then minimizing this upper bound (minimization step).

The issue with this majorization is that it is very coarse: in dimension 2, take $f(x) = x_1^2 + 0.001x^2$. For this function, $L = 1$ and $\mu = 0.001$: we get a linear convergence rate for gradient descent, but the constant is very poor (see also Figure 2), because the upper bounding parabolas that we construct have Hessian $L \text{Id}$, and $\mu \ll L$. In practice, this leads to very small steps being performed after the first one.

6.1 Newton method

We can try to preserve more information about the curvature, by using a 2nd order Taylor expansion of f – hence the name of “second order methods”:

$$f(y) \approx f(x) + \langle \nabla f(x), y - x \rangle + \frac{1}{2} \langle y - x, \nabla^2 f(x)(y - x) \rangle . \quad (6.1)$$

As in the MM interpretation of GD, we can derive an iterative algorithm for this: set $x = x_k$, and define x_{k+1} as the minimizer of the approximation. This yields

$$x_{k+1} = x_k - [\nabla^2 f(x_k)]^{-1} \nabla f(x_k) . \quad (6.2)$$

Compared to gradient descent, the scalar stepsize has been replaced with the application of a PSD matrix, that better captures the curvature of f .

Let us illustrate the possible behaviors for Newton’s method on a simple 1-dimensional convex function, $x \mapsto \sqrt{1+x^2}$. This function is \mathcal{C}^2 , with $f'(x) = \frac{x}{\sqrt{1+x^2}}$ and $f''(x) = \frac{1}{(1+x^2)^{3/2}}$, so Newton’s method iterates are given by

$$x_{k+1} = x_k - x_k(1+x_k^2) = -x_k^3 . \quad (6.3)$$

There are therefore, depending on $|x_0|$, three possible behaviors:

- extremely fast convergence towards the minimizer if $|x_0| < 1$
- oscillations between 1 and -1 if $|x_0| = 1$
- extremely fast divergence if $|x_0| > 1$

This examples illustrates a crucial fact about the vanilla Newton method: it only has *local* convergence properties – it converges if initialized close enough to the minimizer. But when it does, it enjoys a *superlinear* convergence rate. Let us formalize this and provide a proof.

Proposition 6.1. *If $x^* = \operatorname{argmin} f(x)$, f has M -Lipschitz Hessian, $\nabla^2 f(x^*) \succeq \mu \text{Id}_n$ ($\mu > 0$), and $\|x_0 - x^*\| \leq \frac{\mu}{2M}$, then:*

$$\|x_{k+1} - x^*\| \leq \frac{M}{\mu} \|x_k - x^*\|^2 . \quad (6.4)$$

It is hard to imagine a better rate than this one: the sequence $u_k = \frac{M}{\mu} \|x_k - x^*\|$ satisfies $u_{k+1} \leq u_k^2 \leq u_0^{2^k}$. If $u_0 < 1$, the number of iterations to reach $u_k < \varepsilon$ is of the order $\log(\log \varepsilon)$, that is, in practice, basically a constant!

Proof. By the fundamental theorem of calculus, for all x, h in \mathbb{R}^d :

$$\nabla f(x+h) - \nabla f(x) = \int_0^1 \nabla^2 f(x+th)h \, dt \quad (6.5)$$

Hence

$$\nabla f(x_k) = \int_0^1 \nabla^2 f(x^* + t(x_k - x^*))(x_k - x^*) \, dt, \quad (6.6)$$

and thus

$$x_{k+1} - x^* = x_k - x^* - [\nabla^2 f(x_k)]^{-1} \int_0^1 \nabla^2 f(x^* + t(x_k - x^*))(x_k - x^*) \, dt \quad (6.7)$$

$$= [\nabla^2 f(x_k)]^{-1} \int_0^1 \left(\nabla^2 f(x_k) - \nabla^2 f(x^* + t(x_k - x^*)) \right) (x_k - x^*) \, dt. \quad (6.8)$$

Let's control the operator norm of the matrix being integrated, using M -Lipschitzness of the Hessian

$$\int_0^1 \|\nabla^2 f(x_k) - \nabla^2 f(x^* + t(x_k - x^*))\| \, dt \leq \int_0^1 (1-t)M\|x_k - x^*\| \, dt = \frac{M}{2}\|x_k - x^*\| \quad (6.9)$$

Finally, let's bound the operator norm of $[\nabla^2 f(x_k)]^{-1}$:

$$\nabla^2 f(x_k) \succeq \nabla^2 f(x^*) - M\|x_k - x^*\| \text{Id}_n \succeq (\mu - M \underbrace{\|x_k - x^*\|}_{\leq \frac{\mu}{2M} \text{ (induction)}}) \text{Id}_n \succeq \frac{\mu}{2} \text{Id}_n \quad (6.10)$$

□

A few more comments: it is a descent method, it converges in one step for parabolas (well, it makes sense, because the 2nd order approximation of f that we minimize is f itself). But it is not the “one algorithm to rule them all”:

- it may diverge if not started close enough to optimum
- the Hessian may not be invertible
- the Hessian is costly to compute
- the Hessian is even more costly to invert ($\mathcal{O}(d^3)$)

The purpose of Quasi-Newton methods is to replace the inverse Hessian in [Equation \(6.2\)](#) by less costly approximations, in order to retain the very fast convergence rate while alleviating the computational burden.

6.2 Quasi Newton methods

Algorithms in this section will have the following form:

$$\begin{cases} d_k = -B_k g_k \\ x_{k+1} = x_k + \rho_k d_k \end{cases} \quad \text{or} \quad \begin{cases} d_k = -H_k^{-1} g_k \\ x_{k+1} = x_k + \rho_k d_k \end{cases} \quad (6.11)$$

where $g_k = \nabla f(x_k)$, d_k is to be thought of as a descent direction, H_k (resp. B_k) is an approximation of the Hessian (resp. the inverse Hessian) of f at x_k , and ρ_k is a stepsize.

How to build such approximations? It is reasonable to impose that approximations should satisfy the *secant condition* or *Quasi-Newton relation*:

$$g_{k+1} - g_k = H_{k+1}(x_{k+1} - x_k) \quad \text{or} \quad x_{k+1} - x_k = B_{k+1}(g_{k+1} - g_k) \quad (6.12)$$

Where does this condition come from? Suppose we have finished iteration $k+1$ and, to compute x_{k+2} , want to construct (then minimize) ϕ_{k+1} , a quadratic approximation of f having H_{k+1} as Hessian:

$$\phi_{k+1}(x) = f(x_{k+1}) + \langle g_{k+1}, x - x_{k+1} \rangle + \frac{1}{2} \langle x - x_{k+1}, H_{k+1}(x - x_{k+1}) \rangle . \quad (6.13)$$

This quadratic approximation of f is exact at x_{k+1} ($\phi(x_{k+1}) = f(x_{k+1})$), and tangent at x_{k+1} too ($\nabla \phi(x_{k+1}) = \nabla f(x_{k+1})$). To impose the secant condition is simply to impose that f and ϕ_{k+1} are also tangent at the previous iterate, x_k .

6.2.1 The Broyden/SR1 formula

The approximations in this section are updated iteratively, using rank one or rank 2 corrections from B_k/H_k to B_{k+1}/H_{k+1} .

Following the notation of Nocedal and Wright, we write

$$y_k = g_{k+1} - g_k , \quad (6.14)$$

$$s_k = x_{k+1} - x_k . \quad (6.15)$$

Consider an update of the form $H_{k+1} = H_k + \sigma v v^\top$. How should we set σ and v so that it satisfies the secant condition (6.12)? We want:

$$y_k = H_{k+1} s_k \quad (6.16)$$

$$= H_k s_k + \sigma v v^\top s_k \quad (6.17)$$

$$= H_k s_k + (\sigma v^\top s_k) v \quad (\text{rotation trick!}) , \quad (6.18)$$

This does not give us v , but it gives its direction: v is proportional to $y_k - H_k s_k$. So we know⁵ that there exists a scalar δ such that $v = \delta(y_k - H_k s_k)$, and we can plug it back in Equation (6.17) to solve for δ and obtain that

$$H_{k+1} = H_k + \frac{1}{s_k^\top (y_k - H_k s_k)} (y_k - H_k s_k)(y_k - H_k s_k)^\top . \quad (6.19)$$

⁵what if $v^\top s_k = 0$?

satisfies the secant condition. This is called the Broyden or SR1 formula.

This approximation removes the cost of computing the Hessian, and also its cost of storing (we just store the scalars $s_k^\top(y_k - H_k s_k)$ and vectors $(y_k - H_k s_k)$). But at first sight we still have the $\mathcal{O}(d^3)$ to compute the update (linear system solving). It turns out that we can also get rid of it, using Sherman-Woodbury-Morrison formula ([Exercise 6.3](#)):

$$(H + uv^\top)^{-1} = H^{-1} - \frac{1}{1 + v^\top H^{-1}u} (H^{-1}u)(v^\top H^{-1}) . \quad (6.20)$$

This formula means that an update of rank 1 in H_k leads to an update of rank 1 in H_k^{-1} . If we start at a simple H_0 , say Id, then we only need to do, at each iteration, rank 1 updates that are cheap to store and apply ([Exercise 6.2](#))!

6.2.2 DFP and BFGS

6.3 Exercises

Exercise 6.2. ☹ Show that any $n \times n$ matrix of rank 1 writes xy^\top with $x, y \in \mathbb{R}^n$. Show that the cost of multiplying by xy^\top is $2n$ instead of the usual n^2 .

Exercise 6.3 (Sherman-Woodbury-Morrison inverse formula). ☹☹ Let $n, k \in \mathbb{N}$. Let $U \in \mathbb{R}^{n \times k}$ and $V \in \mathbb{R}^{k \times n}$. Show that

$$(\text{Id} + UV)^{-1} = \text{Id} - U(\text{Id} + VU)^{-1}V . \quad (6.21)$$

Show that for any invertible A ,

$$(A + UV)^{-1} = A^{-1} - A^{-1}U(\text{Id} + VA^{-1}U)^{-1}VA^{-1} . \quad (6.22)$$